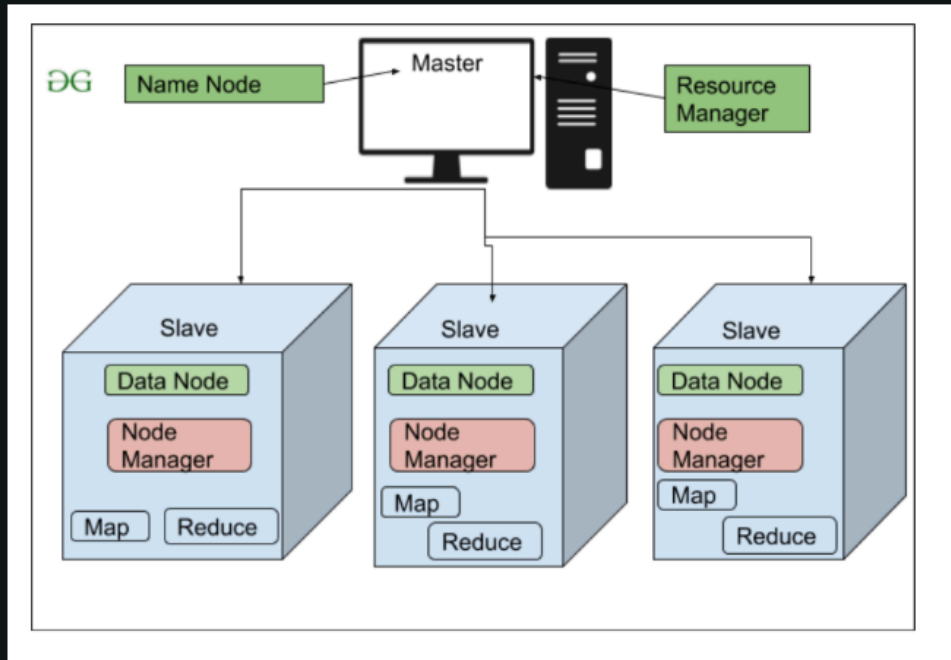


# High Level Architecture Of Hadoop

## High Level Architecture Of Hadoop



### 1. Why is Hadoop slow in processing data?

- Hadoop's MapReduce framework is **disk-based**, meaning it writes intermediate results to disk between each stage (Map → Shuffle → Reduce).
- This frequent **disk I/O** causes **high latency** and makes Hadoop **slow**, especially for iterative and real-time tasks.

### 2. How does Apache Spark improve processing speed?

- Spark uses **in-memory computation**, meaning it stores and processes data in **RAM instead of disk**.
- This reduces disk I/O and makes Spark **up to 100x faster** than Hadoop for certain workloads.

---

### 3. Why is Hadoop's programming model complex?

- Writing MapReduce jobs requires **manual handling of data flow**, which results in **long, complex code**.
- Developers need to write multiple jobs and manually manage dependencies.

#### 4. How does Apache Spark simplify programming?

- Spark provides **high-level APIs** in **Python (PySpark), Scala, Java, and R**, making it much easier to use.
  - It supports **SQL (Spark SQL), streaming (Spark Streaming), and machine learning (MLlib)**, reducing the need for complex code.
- 

#### 5. Why is Hadoop inefficient for iterative processing?

- Many tasks, like **machine learning and graph processing**, require running multiple iterations over the same data.
- Hadoop **reloads data from disk in each iteration**, making it inefficient and slow.

#### 6. How does Spark handle iterative processing better?

- Spark introduces **Resilient Distributed Datasets (RDDs)**, which keep intermediate results **in memory**.
  - This eliminates redundant disk reads and makes iterative tasks **significantly faster**.
- 

#### 7. Can Hadoop handle real-time data processing?

- No, Hadoop's batch processing nature makes it **unsuitable for real-time workloads**.
- It processes data in large batches, causing delays in real-time decision-making.

#### 8. How does Spark enable real-time data processing?

- Spark Streaming processes data in **real-time** using **micro-batches**, allowing low-latency data processing.
  - This makes Spark ideal for real-time applications like fraud detection and IoT analytics.
- 

#### 9. Why does Hadoop have high latency?

- Due to frequent disk writes, Hadoop's MapReduce jobs experience **high latency** and slow execution times.

10. How does Spark reduce latency?

- Spark processes data **in-memory** and uses **optimized execution plans (DAG – Directed Acyclic Graph)**.
  - This significantly reduces latency and speeds up execution.
- 

11. Why is debugging and maintaining Hadoop difficult?

- Debugging MapReduce jobs is challenging because of the **distributed execution** and **complex logs**.
- Performance tuning requires **deep technical expertise**.

12. How does Spark improve debugging and maintenance?

- Spark’s high-level APIs and **interactive shell** (PySpark, Spark Shell) make it easier to debug and optimize jobs.
  - The DAG execution model provides **better job monitoring and fault tolerance**.
- 

13. What makes Spark a better choice than Hadoop for modern applications?

- Spark offers a **unified framework** for **batch processing, real-time streaming, machine learning, and SQL queries**.
  - Unlike Hadoop, which needs **multiple tools** (e.g., Hadoop + Storm + Hive), Spark handles everything in a **single platform**.
- 

Summary Table: Hadoop vs. Spark

Feature	Hadoop (MapReduce)	Apache Spark
Processing Speed	Slow (disk-based)	Fast (in-memory)
Programming Complexity	High (manual MapReduce code)	Easy (high-level APIs)
Real-time Processing	No (batch-oriented)	Yes (Spark Streaming)
Iterative Processing	Inefficient (reloads from disk)	Efficient (keeps data in memory)
Latency	High	Low

## Use Cases

Batch processing

Batch + Streaming + ML +  
SQL

Apache Spark can be used in two major data architecture solutions:

- 1 Data Lake on Hadoop
- 2 Lakehouse on Cloud

Let's explore both in detail:

---

## 1. Data Lake on Hadoop (HDFS-based Data Lake)

### What is it?

- A **data lake** is a centralized repository that stores **raw, unstructured, semi-structured, and structured** data.
- It is built **on Hadoop's HDFS (Hadoop Distributed File System)** and uses **Apache Spark** for processing.

### Key Components

- ✓ **HDFS** – Stores raw data in a distributed manner.
- ✓ **YARN** – Manages cluster resources.
- ✓ **Apache Spark** – Processes large-scale data stored in HDFS.
- ✓ **Hive, Presto, Impala** – SQL-based querying on data lake.

### Use Cases

- ✓ Batch processing & ETL
- ✓ Big data analytics
- ✓ Storing massive raw datasets (logs, IoT, social media)

### Example Workflow

- 1 Raw data is ingested into **HDFS** (from databases, IoT devices, logs).
- 2 **Apache Spark** processes and transforms the data.

- ③ **Hive or Presto** enables SQL-based querying.
- ④ Data is used for reporting, analytics, or machine learning.

## Advantages

- ✓ Cost-effective storage for large datasets
- ✓ Scales horizontally with Hadoop clusters
- ✓ Can handle any data format (structured, semi-structured, unstructured)

## Disadvantages

- ✗ High latency (not great for real-time analytics)
  - ✗ Data quality issues due to schema-on-read
  - ✗ Complex to manage (requires Hadoop expertise)
- 

## 2. Lakehouse on Cloud (Modern Data Lakehouse)

### What is it?

- A **Lakehouse** combines the **best of Data Lakes and Data Warehouses** by providing:
  - The **scalability** of a data lake
  - The **structure & ACID transactions** of a data warehouse
- Built on **cloud storage** (S3, Azure Blob, Google Cloud Storage) with **Apache Spark** for fast analytics.

### Key Components

- ✓ **Cloud Object Storage** – (S3, Azure Blob, GCS) stores data in an open format (Parquet, Delta Lake).
- ✓ **Delta Lake / Iceberg / Hudi** – Adds **ACID transactions** and schema enforcement.
- ✓ **Apache Spark** – Handles ETL, ML, and analytics.
- ✓ **Databricks / Snowflake** – Provides high-performance query engines.

### Use Cases

- ✓ Real-time data processing
- ✓ Machine learning & AI
- ✓ Enterprise analytics & BI

### Example Workflow

- ① **Raw data lands in cloud storage** (S3, Azure Blob, GCS).
- ② **Delta Lake/Iceberg/Hudi** manages transactions and schema.

- 3 Apache Spark processes & transforms data efficiently.
- 4 BI tools (Tableau, Power BI) or ML models consume the clean data.

Advantages

- ✓ ACID transactions ensure data integrity.
- ✓ Faster than Hadoop (optimized for cloud storage & in-memory processing).
- ✓ Real-time & interactive queries are possible.
- ✓ Easier to manage than Hadoop-based lakes.

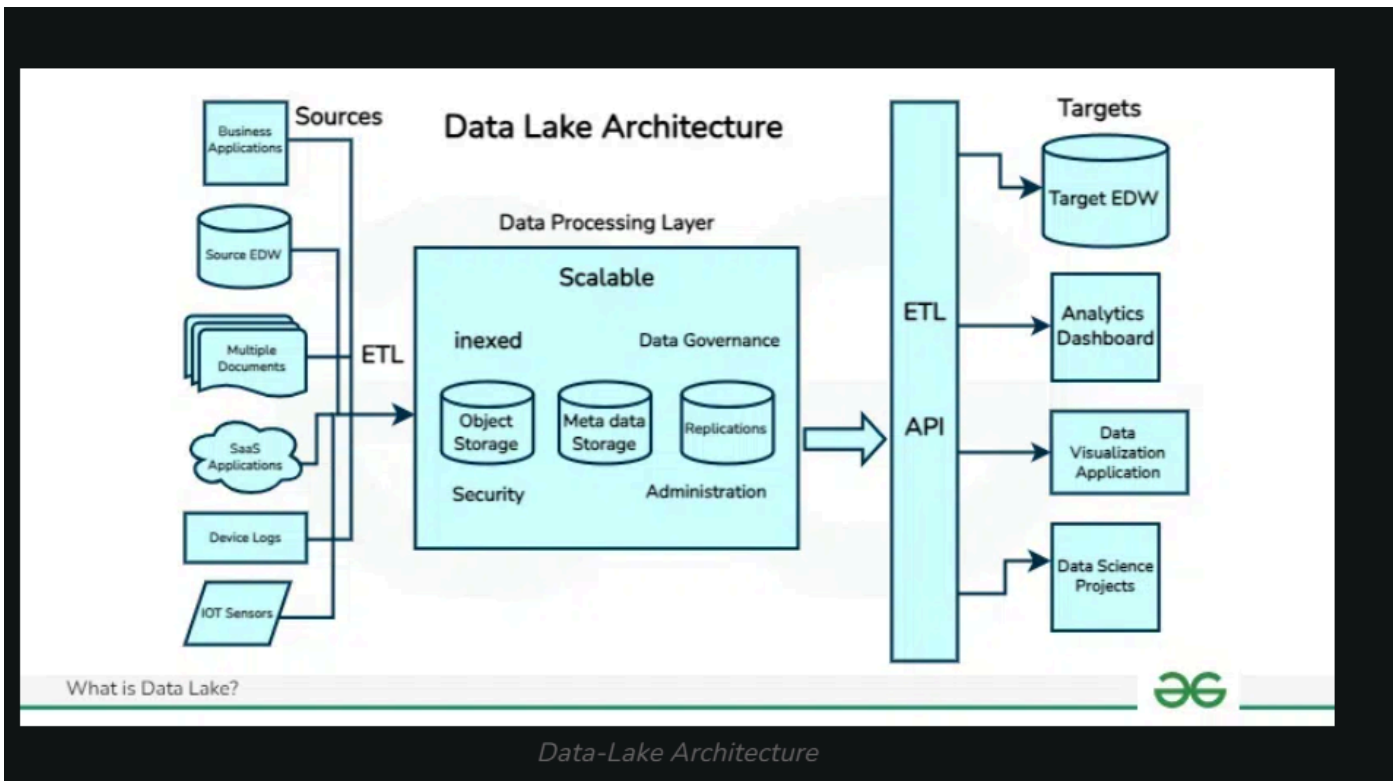
Disadvantages

- ✗ Cloud costs can increase with high storage/compute usage.
- ✗ Vendor lock-in (Databricks, Snowflake, etc.).

Comparison: Data Lake (Hadoop) vs. Lakehouse (Cloud)

Feature	Data Lake (HDFS)	Lakehouse (Cloud)
Storage	HDFS (on-prem)	S3, Azure Blob, GCS (cloud)
Processing	Apache Spark, Hive, Presto	Apache Spark, Delta Lake, Iceberg
Schema	Schema-on-read	Schema enforcement (ACID)
Query Speed	Slower (batch-oriented)	Faster (real-time & interactive)
Cost	Cheaper storage, high infra cost	Pay-as-you-go, scalable
Management	Complex (requires Hadoop expertise)	Easier (cloud-managed services)

# Data Lake Architecture Diagram



## 1. Storage Layer (Centralized Repository)

- Stores raw and processed data in a distributed file system or cloud storage.
- **Storage options:**
  - **HDFS** (for on-premises data lakes)
  - **Cloud Object Storage** (Amazon S3, Azure Blob, Google Cloud Storage)
- Supports various formats: **CSV, JSON, Avro, Parquet, ORC**.

## 3. Processing Layer (Data Transformation & Analytics)

- Transforms raw data into usable formats using **batch or real-time processing**.
- **Batch Processing** (ETL, data cleaning): Apache Spark, Hive, Presto
- **Real-time Processing** (Streaming data): Apache Flink, Spark Streaming, Kafka Streams

## ○ Data Lake Architecture working

